

Prompt engineering: progress

Prompt engineering: progress

My hands-on log for learning prompt engineering. Goal: be able to take a vague request and turn it into a precise, tested prompt, and know which lever (clarity, examples, format, constraints, reasoning, grounding) to reach for.

Status: [] not started, [~] in progress, [x] done

Module	Technique	Status
01	Anatomy + clarity/specificity	[]
02	Few-shot (in-context examples)	[]
03	Structured output (JSON)	[]
04	Roles and constraints	[]
05	Chain-of-thought	[]
06	Grounding / anti-hallucination	[]
07	Iterate and evaluate	[]
08	Pitfalls and injection (reading)	[]

Setup

- ☐ Ollama running with qwen2.5:3b-instruct (ollama serve)

Per module

- ☐ 01: ran the demo, saw weak vs strong, then wrote my own specific prompt
- ☐ 02: saw few-shot make the output consistent; tried changing the examples
- ☐ 03: got valid JSON out and parsed it; broke it on purpose to see why format rules matter
- ☐ 04: controlled tone/length with a system prompt + constraints
- ☐ 05: saw step-by-step fix a multi-step problem; noticed the verbosity cost
- ☐ 06: made the model abstain on an unanswerable question (grounding)
- ☐ 07: ran the A/B harness and picked the better prompt by score, not by feel
- ☐ 08: read the pitfalls + injection notes

What I learned / surprised me
